

Notebook Overview: Tuning for HAN-GRU Model #2

We aimed to enhance the model's performance by modifying specific hyperparameters. We increased the attention size from 100 to 200 to enable a more comprehensive capture of intricate lyrical relationships. However, this adjustment posed a potential risk of overfitting. To counter this, we raised the dropout rate from 0.3 to 0.4, mitigating the risk of overfitting by introducing more regularization. Additionally, we reduced the hidden size to encourage the model to generalize better on unseen data.

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
from tqdm import tqdm
import re

import torch
from torch import nn
import torch.optim as optim
from torch.nn.utils.rnn import pad_sequence
from torch.utils.data import Dataset, DataLoader, random_split, RandomSampler, SequentialSampler

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.feature_extraction.text import CountVectorizer
from torch.utils.data import DataLoader, TensorDataset
from sklearn.decomposition import TruncatedSVD

from torchtext.data import get_tokenizer
from collections import Counter
from torchtext.vocab import Vocab, build_vocab_from_iterator
```

```
In [2]: df = pd.read_csv("data/lyrics_cleaned.csv")
```

```
In [3]: tokenizer = get_tokenizer('basic_english')
counter = Counter()
for line in tqdm(df['lyrics']):
    counter.update(tokenizer(line))
```

100%|██████████| 218162/218162 [00:27<00:00, 7959.63it/s]

```
In [4]: # Create vocabulary using build_vocab_from_iterator
vocab = build_vocab_from_iterator([tokenizer(line) for line in df['lyrics']],
                                specials=['<unk>', '<pad>'], min_freq=1)
```

```
In [5]: label_encoder = LabelEncoder()
indexed_data = [torch.tensor([vocab[token] for token in tokenizer(line)])
                for line in df['lyrics']]
```

```
# Include padding for same shape size
max_seq_length = max(len(seq) for seq in indexed_data)
```

```
In [6]: padded_data = pad_sequence(indexed_data, batch_first=True, padding_value=vocab['<pad>'])
indexed_labels = torch.tensor(label_encoder.fit_transform(df['genre']))
```

```
In [7]: class LyricsDataset(Dataset):
        def __init__(self, lyrics, genre):
            self.lyrics = lyrics
            self.genre = genre

        def __len__(self):
            return len(self.genre)

        def __getitem__(self, idx):
            return self.lyrics[idx], self.genre[idx]

dataset = LyricsDataset(padded_data, indexed_labels)

# Split into training and validation sets
train_size = int(0.8 * len(dataset))
val_size = len(dataset) - train_size

train_dataset, val_dataset = random_split(dataset, [train_size, val_size])

print('{:>5,} training samples'.format(train_size))
print('{:>5,} validation samples'.format(val_size))
```

174,529 training samples
43,633 validation samples

```
In [8]: batch_size=32

train_dataloader = DataLoader(
    train_dataset,
    sampler = RandomSampler(train_dataset),
    batch_size = batch_size
)

validation_dataloader = DataLoader(
    val_dataset,
    sampler = SequentialSampler(val_dataset),
    batch_size = batch_size
)
```

Model

```
In [10]: import torch
import torch.nn as nn

class AttLayer(nn.Module):
    def __init__(self, input_size, hidden_dim):
        super(AttLayer, self).__init__()
        self.hidden_dim = hidden_dim
        self.W = nn.Parameter(torch.randn(input_size, hidden_dim))
        self.bw = nn.Parameter(torch.zeros(hidden_dim))
        self.uw = nn.Parameter(torch.randn(hidden_dim))

    def forward(self, x):
        batch_size, num_words, hidden_size = x.size()
        x_resaped = x.reshape(-1, hidden_size)

        ui = torch.tanh(torch.matmul(x_resaped, self.W) + self.bw)
        intermed = torch.sum(self.uw * ui, dim=1)

        intermed = intermed.view(batch_size, num_words)
        weights = torch.softmax(intermed, dim=-1)
        weights = weights.unsqueeze(-1)

        weighted_input = x * weights
        return torch.sum(weighted_input, dim=1)
```

```

class HAN_GRU(nn.Module):
    def __init__(self, num_words, embedding_vector_length, hidden_size, attention_size, max_words_per_line):
        super(HAN_GRU, self).__init__()

        self.word_embedding = nn.Embedding(num_words, embedding_vector_length)
        self.word_gru = nn.GRU(embedding_vector_length, hidden_size, batch_first=True, bidirectional=True)
        self.word_attention = AttLayer(hidden_size * 2, attention_size)

        self.sentence_gru = nn.GRU(hidden_size * 2, hidden_size, batch_first=True, bidirectional=True)
        self.sentence_attention = AttLayer(hidden_size * 2, attention_size)

        self.max_words_per_line = max_words_per_line
        self.max_num_lines = max_num_lines

        self.dropout = nn.Dropout(0.4)
        self.fc = nn.Linear(hidden_size * 2, output_size)

    def forward(self, inputs):
        word_embedded = self.word_embedding(inputs)

        word_output, _ = self.word_gru(word_embedded)
        word_attention_output = self.word_attention(word_output)

        batch_size = word_attention_output.size(0)
        sentence_input = word_attention_output.view(batch_size, -1, word_attention_output.size(-1))

        sentence_output, _ = self.sentence_gru(sentence_input)
        sentence_attention_output = self.sentence_attention(sentence_output)

        document_output = sentence_attention_output.view(batch_size, -1)
        output = self.fc(self.dropout(document_output))
        return output

```

Parameters changed:

- Original attention size was 100; changed to 200
- why? to try capture more complex relationships of the lyrics with the risk of potentially overfitting

- Dropout rate from 0.3 --> 0.4
- why? improve generalization and prevent overfitting
- hidden size from 128 --> 96
- why? Smaller hidden sizes may lead to a more regularized model that generalizes better to unseen data. It prevents the model from memorizing the training data, forcing it to learn more abstract and useful representations. --> try obtain patterns from lyrics

```
In [11]: # Calculate max words per line and max number of lines
max_words_per_line = df['lyrics'].apply(lambda x: len(x.split())).max()
max_num_lines = df['lyrics'].apply(lambda x: len(x.split('\n'))).max()
print(f"Max Words Per Line: {max_words_per_line}")
print(f"Max Number of Lines: {max_num_lines}")

attention_size = 200
hidden_size = 96
vocab_size = len(vocab)
embedding_dim = 512
output_size = len(df['genre'].unique())
```

Max Words Per Line: 6232

Max Number of Lines: 759

```
In [12]: num_epochs = 10
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

model = HAN_GRU(vocab_size, embedding_dim, hidden_size, attention_size,
                 max_words_per_line, max_num_lines, output_size).to(device)
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
criterion = nn.CrossEntropyLoss()

print(model)
```

```

HAN_GRU(
    (word_embedding): Embedding(245576, 512)
    (word_gru): GRU(512, 96, batch_first=True, bidirectional=True)
    (word_attention): AttLayer()
    (sentence_gru): GRU(192, 96, batch_first=True, bidirectional=True)
    (sentence_attention): AttLayer()
    (dropout): Dropout(p=0.4, inplace=False)
    (fc): Linear(in_features=192, out_features=11, bias=True)
)

```

```

In [14]: losses, accuracies = [], []

for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0
    correct_predictions = 0
    total_predictions = 0

    for inputs, labels in tqdm(train_dataloader):
        inputs, labels = inputs.to(device), labels.to(device)

        optimizer.zero_grad()

        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()

        _, predicted = torch.max(outputs, 1)
        total_predictions += labels.size(0)
        correct_predictions += (predicted == labels).sum().item()

    epoch_loss = running_loss / len(train_dataloader)
    epoch_accuracy = (correct_predictions / total_predictions) * 100
    losses.append(epoch_loss)
    accuracies.append(epoch_accuracy)

    print(f"Epoch [{epoch + 1}/{num_epochs}] Train Loss: {epoch_loss:.4f} Train Accuracy: {epoch_accuracy:.4f}")

```

```
100%|██████████| 5455/5455 [18:25<00:00, 4.94it/s]
```

```

Epoch [1/10] Train Loss: 1.5164 Train Accuracy: 52.26%
100%|██████████| 5455/5455 [18:23<00:00, 4.94it/s]
Epoch [2/10] Train Loss: 1.4642 Train Accuracy: 53.83%
100%|██████████| 5455/5455 [18:22<00:00, 4.95it/s]
Epoch [3/10] Train Loss: 1.4317 Train Accuracy: 54.48%
100%|██████████| 5455/5455 [18:22<00:00, 4.95it/s]
Epoch [4/10] Train Loss: 1.3997 Train Accuracy: 55.45%
100%|██████████| 5455/5455 [18:21<00:00, 4.95it/s]
Epoch [5/10] Train Loss: 1.3794 Train Accuracy: 56.09%
100%|██████████| 5455/5455 [18:23<00:00, 4.94it/s]
Epoch [6/10] Train Loss: 1.3563 Train Accuracy: 56.70%
100%|██████████| 5455/5455 [18:24<00:00, 4.94it/s]
Epoch [7/10] Train Loss: 1.3351 Train Accuracy: 57.29%
100%|██████████| 5455/5455 [18:22<00:00, 4.95it/s]
Epoch [8/10] Train Loss: 1.3124 Train Accuracy: 57.95%
100%|██████████| 5455/5455 [18:21<00:00, 4.95it/s]
Epoch [9/10] Train Loss: 1.2893 Train Accuracy: 58.76%
100%|██████████| 5455/5455 [18:25<00:00, 4.93it/s]
Epoch [10/10] Train Loss: 1.2726 Train Accuracy: 59.52%

```

```

In [15]: # Validation loop (optional)
model.eval()
val_running_loss = 0.0
correct = 0
total = 0

with torch.no_grad():
    for inputs, labels in validation_dataloader:
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = model(inputs)
        val_loss = criterion(outputs, labels)
        val_running_loss += val_loss.item()

        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

accuracy = correct / total

```

```
avg_val_loss = val_running_loss / len(validation_dataloader)
print(f'Validation Loss: {avg_val_loss:.4f}, Accuracy: {accuracy * 100:.2f}%')
```

Validation Loss: 1.3865, Accuracy: 55.79%

In []: