

Notebook Overview: Tuning for HAN-GRU Model #3

Our aim was to enhance the complexity of our model by adding an additional layer to our GRU cell.

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
from tqdm import tqdm
import re

import torch
from torch import nn
import torch.optim as optim
from torch.nn.utils.rnn import pad_sequence
from torch.utils.data import Dataset, DataLoader, random_split, RandomSampler, SequentialSampler

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.feature_extraction.text import CountVectorizer
from torch.utils.data import DataLoader, TensorDataset
from sklearn.decomposition import TruncatedSVD

from torchtext.data import get_tokenizer
from collections import Counter
from torchtext.vocab import Vocab, build_vocab_from_iterator
```

```
In [2]: df = pd.read_csv("data/lyrics_cleaned.csv")
```

```
In [3]: tokenizer = get_tokenizer('basic_english')
counter = Counter()
for line in tqdm(df['lyrics']):
    counter.update(tokenizer(line))
```

```
100%|██████████| 218162/218162 [00:27<00:00, 7920.55it/s]
```

```
In [4]: # Create vocabulary using build_vocab_from_iterator
vocab = build_vocab_from_iterator([tokenizer(line) for line in df['lyrics']],
```

```
specials=['<unk>', '<pad>'], min_freq=1)
```

```
In [5]: label_encoder = LabelEncoder()
indexed_data = [torch.tensor([vocab[token] for token in tokenizer(line)])
                for line in df['lyrics']]

# Include padding for same shape size
max_seq_length = max(len(seq) for seq in indexed_data)
```

```
In [6]: padded_data = pad_sequence(indexed_data, batch_first=True, padding_value=vocab['<pad>'])
indexed_labels = torch.tensor(label_encoder.fit_transform(df['genre']))
```

```
In [7]: class LyricsDataset(Dataset):
        def __init__(self, lyrics, genre):
            self.lyrics = lyrics
            self.genre = genre

        def __len__(self):
            return len(self.genre)

        def __getitem__(self, idx):
            return self.lyrics[idx], self.genre[idx]

dataset = LyricsDataset(padded_data, indexed_labels)

# Split into training and validation sets
train_size = int(0.8 * len(dataset))
val_size = len(dataset) - train_size

train_dataset, val_dataset = random_split(dataset, [train_size, val_size])

print('{:>5,} training samples'.format(train_size))
print('{:>5,} validation samples'.format(val_size))
```

```
174,529 training samples
43,633 validation samples
```

```
In [8]: batch_size=32

train_dataloader = DataLoader(
    train_dataset,
```

```

        sampler = RandomSampler(train_dataset),
        batch_size = batch_size
    )

validation_dataloader = DataLoader(
    val_dataset,
    sampler = SequentialSampler(val_dataset),
    batch_size = batch_size
)

```

Model

```

In [9]: import torch
import torch.nn as nn

class AttLayer(nn.Module):
    def __init__(self, input_size, hidden_dim):
        super(AttLayer, self).__init__()
        self.hidden_dim = hidden_dim
        self.W = nn.Parameter(torch.randn(input_size, hidden_dim))
        self.bw = nn.Parameter(torch.zeros(hidden_dim))
        self.uw = nn.Parameter(torch.randn(hidden_dim))

    def forward(self, x):
        batch_size, num_words, hidden_size = x.size()
        x_resaped = x.reshape(-1, hidden_size)

        ui = torch.tanh(torch.matmul(x_resaped, self.W) + self.bw)
        intermed = torch.sum(self.uw * ui, dim=1)

        intermed = intermed.view(batch_size, num_words)
        weights = torch.softmax(intermed, dim=-1)
        weights = weights.unsqueeze(-1)

        weighted_input = x * weights
        return torch.sum(weighted_input, dim=1)

class HAN_GRU(nn.Module):
    def __init__(self, num_words, embedding_vector_length, hidden_size, attention_size, max_words_per_line
        super(HAN_GRU, self).__init__()

```

```

self.word_embedding = nn.Embedding(num_words, embedding_vector_length)
self.word_gru = nn.GRU(embedding_vector_length, hidden_size, batch_first=True, bidirectional=True, num_
self.word_attention = AttLayer(hidden_size * 2, attention_size)

self.sentence_gru = nn.GRU(hidden_size * 2, hidden_size, batch_first=True, bidirectional=True, num_
self.sentence_attention = AttLayer(hidden_size * 2, attention_size)

self.max_words_per_line = max_words_per_line
self.max_num_lines = max_num_lines

self.dropout = nn.Dropout(0.4)
self.fc = nn.Linear(hidden_size * 2, output_size)

def forward(self, inputs):
    word_embedded = self.word_embedding(inputs)

    word_output, _ = self.word_gru(word_embedded)
    word_attention_output = self.word_attention(word_output)

    batch_size = word_attention_output.size(0)
    sentence_input = word_attention_output.view(batch_size, -1, word_attention_output.size(-1))

    sentence_output, _ = self.sentence_gru(sentence_input)
    sentence_attention_output = self.sentence_attention(sentence_output)

    document_output = sentence_attention_output.view(batch_size, -1)
    output = self.fc(self.dropout(document_output))
    return output

```

Parameters changed:

- added `num_layers` to GRU cells = 2 (both word and sentence)
- why? adding more complexity in our models
- attention size --> 64
- embedding dim --> 1024 from ver 2

```

In [10]: # Calculate max words per line and max number of lines
max_words_per_line = df['lyrics'].apply(lambda x: len(x.split())).max()

```

```

max_num_lines = df['lyrics'].apply(lambda x: len(x.split('\n'))).max()
print(f"Max Words Per Line: {max_words_per_line}")
print(f"Max Number of Lines: {max_num_lines}")

attention_size = 64
hidden_size = 128
vocab_size = len(vocab)
embedding_dim = 256
output_size = len(df['genre'].unique())

```

Max Words Per Line: 6232

Max Number of Lines: 759

```

In [11]: num_epochs = 10
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

model = HAN_GRU(vocab_size, embedding_dim, hidden_size, attention_size,
                max_words_per_line, max_num_lines, output_size).to(device)
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
criterion = nn.CrossEntropyLoss()

print(model)

```

```

HAN_GRU(
  (word_embedding): Embedding(245576, 256)
  (word_gru): GRU(256, 128, num_layers=2, batch_first=True, bidirectional=True)
  (word_attention): AttLayer()
  (sentence_gru): GRU(256, 128, num_layers=2, batch_first=True, bidirectional=True)
  (sentence_attention): AttLayer()
  (dropout): Dropout(p=0.4, inplace=False)
  (fc): Linear(in_features=256, out_features=11, bias=True)
)

```

```

In [12]: losses, accuracies = [], []

for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0
    correct_predictions = 0
    total_predictions = 0

    for inputs, labels in tqdm(train_data_loader):
        inputs, labels = inputs.to(device), labels.to(device)

```

```

optimizer.zero_grad()

outputs = model(inputs)
loss = criterion(outputs, labels)
loss.backward()
optimizer.step()

running_loss += loss.item()

_, predicted = torch.max(outputs, 1)
total_predictions += labels.size(0)
correct_predictions += (predicted == labels).sum().item()

epoch_loss = running_loss / len(train_dataloader)
epoch_accuracy = (correct_predictions / total_predictions) * 100
losses.append(epoch_loss)
accuracies.append(epoch_accuracy)

print(f"Epoch [{epoch + 1}/{num_epochs}] Train Loss: {epoch_loss:.4f} Train Accuracy: {epoch_accuracy:

```

```

100%|██████████| 5455/5455 [24:46<00:00, 3.67it/s]
Epoch [1/10] Train Loss: 1.3958 Train Accuracy: 55.79%
100%|██████████| 5455/5455 [24:51<00:00, 3.66it/s]
Epoch [2/10] Train Loss: 1.2307 Train Accuracy: 60.85%
100%|██████████| 5455/5455 [24:52<00:00, 3.66it/s]
Epoch [3/10] Train Loss: 1.1390 Train Accuracy: 64.05%
100%|██████████| 5455/5455 [24:48<00:00, 3.66it/s]
Epoch [4/10] Train Loss: 1.0437 Train Accuracy: 67.31%
100%|██████████| 5455/5455 [24:46<00:00, 3.67it/s]
Epoch [5/10] Train Loss: 0.9447 Train Accuracy: 70.83%
100%|██████████| 5455/5455 [24:51<00:00, 3.66it/s]
Epoch [6/10] Train Loss: 0.8457 Train Accuracy: 74.04%
100%|██████████| 5455/5455 [24:49<00:00, 3.66it/s]
Epoch [7/10] Train Loss: 0.7566 Train Accuracy: 76.87%
100%|██████████| 5455/5455 [24:53<00:00, 3.65it/s]
Epoch [8/10] Train Loss: 0.6832 Train Accuracy: 79.14%
100%|██████████| 5455/5455 [24:53<00:00, 3.65it/s]
Epoch [9/10] Train Loss: 0.6211 Train Accuracy: 80.84%

```

```
100%|██████████| 5455/5455 [24:52<00:00, 3.65it/s]  
Epoch [10/10] Train Loss: 0.5707 Train Accuracy: 82.32%
```

```
In [13]: # Validation loop (optional)  
model.eval()  
val_running_loss = 0.0  
correct = 0  
total = 0  
  
with torch.no_grad():  
    for inputs, labels in validation_dataloader:  
        inputs, labels = inputs.to(device), labels.to(device)  
        outputs = model(inputs)  
        val_loss = criterion(outputs, labels)  
        val_running_loss += val_loss.item()  
  
        _, predicted = torch.max(outputs.data, 1)  
        total += labels.size(0)  
        correct += (predicted == labels).sum().item()  
  
    accuracy = correct / total  
    avg_val_loss = val_running_loss / len(validation_dataloader)  
    print(f'Validation Loss: {avg_val_loss:.4f}, Accuracy: {accuracy * 100:.2f}%')
```

Validation Loss: 1.5832, Accuracy: 59.54%

In []: